

NVIDIA GenAI & LLM Certification Program

Official 5-day NVIDIA certification program on generative AI and large language models, ending in a portfolio capstone

START DATE

6 Sep 2026

DURATION

1 week (5 days)

TOTAL HOURS

40 hrs

PER SEAT

KWD 1,450 / seat

TIMING

8:00 AM – 4:00 PM

[OVERVIEW]

Five days of NVIDIA-aligned certification on large language models and generative AI. Hands-on with NIM, NeMo, and Triton across the full LLM lifecycle. Each attendee leaves with a portfolio project on real data plus the certification voucher.

[WHAT YOU'LL WALK OUT WITH]

- NVIDIA-aligned readiness on GenAI and LLM specialist content
- Hands-on practice across NIM, NeMo, and Triton
- A portfolio project applying NVIDIA's GenAI tooling to a real use case
- Certification voucher in hand at the end of the week

[COHORT DATES]

September 6–10, 2026

[Next cohort](#)

Ready to enrol your team?

enterprise@joincoded.com · coded.kw

NVIDIA GenAI & LLM Certification Program

DAY 1

8 hrs

Fundamentals of Deep Learning

Mechanics of deep learning, training a first computer-vision model, CNNs for image recognition, data augmentation and transfer learning with PyTorch.

DAY 2

8 hrs

Accelerated Data Science

GPU-accelerated data manipulation (cuDF, pandas, Polars, Dask), ML with cuML, XGBoost on GPUs, graph analytics with cuGraph, population-scale data analysis.

DAY 3

8 hrs

Transformer-Based NLP Applications

Transformers as the foundation of modern LLMs, pre-trained LLMs for classification, NER, QA, summarisation, author attribution and chatbot building.

DAY 4

8 hrs

LLM Apps with Prompt Engineering

Prompt engineering principles for production, designing reliable LLM prompts, hands-on LLM application patterns.

DAY 5

8 hrs

Rapid LLM Application Development

End-to-end LLM application development, deployment patterns and integration, capstone build, and certification exam prep.

[WHO THIS IS FOR]

Engineers, ML engineers, and senior IC roles building or operating GenAI systems. Prerequisites: Python proficiency, familiarity with deep-learning fundamentals, comfort with the command line. Maximum cohort size: 35.